

MiniMax H3 • Local Multimodal Video Studio

Open-weights text/image → video + native stereo audio • 100% on a single AMD Radeon GPU (ROCm)

AMD AI DevMaster 2026 • Track 1

ROCm 7.2.4

Radeon PRO W7900 • RDNA3

ComfyUI native nodes

pruned int8 • 48 GB VRAM

Omni-modal in one sampling loop

- **T2V** — text → video + synchronized audio
- **I2V** — + first/last keyframes for shot control
- **R2V** — reference images lock identity / style (<Picture 1>)
- Audio (voice, wind, score) rendered with the video — no TTS pipeline

Generated outputs

shot 1

shot 2

shot 3

System

gradio UI → ComfyUI API → MiniMaxH3 nodes → 2xVAE (video + audio) → mp4

PyTorch ROCm build (HIP) drives ComfyUI — no CUDA paths, reproducible on RDNA2/
RDNA3 ≥24 GB.

Measured on AMD Radeon PRO W7900

Run	Resolution	Duration	Wall time
T2V cold	864x480	~5.2 s	10:42 (load)
T2V warm	608x352	~2.3 s	75 s
T2V warm	864x480	~2.3 s	170 s

Peak VRAM ~27 GB — fits alongside the 27 GB int8 32B text encoder.

Why pruned int8?

- bf16 release is 190 GB+ — cannot fit 48 GB VRAM
- Pruned set removes training-only weights (~21 GB / diffusion)
- int8 (convrot) text encoder runs natively on RDNA

Use cases

Short-form social video • ads & storyboards • e-learning explainers • game cinematics • style-locked R2V character reuse • privacy-safe on-device creation.